

THE JUDGMENT PLATFORM FOR REGULATED ENTERPRISES

We sell speed at AI-native pace.

WisdomTwin is the Judgment Platform for regulated enterprises. With one press you get a live audio huddle with your management team 24/7/365, easier and faster than Zoom, and async wisdom twins trained on decades of executive judgment deliver decisions in minutes with full audit trails.

[WISDOMTWIN.AI](https://wisdomtwin.ai)

[DEMO: WISDOMTWIN.AI/DEMO \(1-4 MIN AUDIO HUDDLES\)](https://wisdomtwin.ai/demo)

[SF-TORONTO + DUBAI-TORONTO = 24/7 BY DESIGN](#)

PROBLEM

Companies don't lose documents when people leave. They lose judgment.

- **Corporate Alzheimer's:** 10,000 Americans turn 65 every day. Judgment walks out the door with no transfer. (US Census)
- Policies record **what** was decided. Never **why**. Principles, exceptions, thresholds, and escalation rules live in heads.
- Regulated teams cannot paste consequential work into a public chatbot, so **meetings remain the bottleneck**.

275

interruptions per employee per day: 117 emails + 153 Teams messages (Microsoft Work Trend Index 2025)

57%

of the day communicating about work vs 43% creating (Microsoft Work Trend Index 2023)

10,000

Americans retire daily. The judgment walks out. (US Census)

Result: decisions that should take minutes wait days.

MISSION

Expand human working memory from 7 items to 1 million, and unleash it 24/7.

Your brain holds about 7 items in working memory. An executive's judgment is bottlenecked into a 40-hour, 9-to-5 week. Pair each role with a wisdom twin holding a million-token context of that role's decisions, and the week changes shape.

7 → 1M

working memory items

40h → 120h

judgment availability per week

3X

speed advantage

Design targets, not measured pilot metrics.

CATEGORY CREATION

Asynchronous management: for the first time, a practical operating model.

- **Zoom didn't invent video calls. Skype did.** Zoom killed friction and owned the category.
- Email, docs, and voice memos tried async leadership. Clunky, incomplete, no judgment layer, no governance.
- **WisdomTwin makes it real. We are the Zoom for management.**

Salesforce said NO SOFTWARE and created SaaS.
We say NO MEETINGS: your twin attends so you don't have to.

COMMERCIAL OUTCOME
We sell speed

CATEGORY
Asynchronous management

PRODUCT
The Judgment Platform

WHY THE BUYER CAN SAY YES
Evidence, permissions, audit, human override

Book meeting



Hunt docs



Wait days

One press



Context live



Decision moves

WHY NOW

Leadership latency has not kept pace with model speed.

The gap is the opportunity.

- **Models reason over decades**, not retrieve chunks. Temporal, conditional, contested judgment is now computable.
- **Regulators accept explainability + human override + audit** as a governance pattern.
- **Private / on-prem / air-gapped inference is practical**: model-agnostic, sovereign-ready.
- **Ontario grants make pilots near-free** for many buyers: procurement friction removed.
- **9-to-5 leadership vs 24/7 models**: asynchronous management closes the gap. Judgment on demand.

Three curves crossed: reasoning + governance acceptance + private deployability. Inevitable, not opportunistic.

PRODUCT

The Judgment Platform: one press to move the decision forward.

1. One-Press Huddle

Press a button: live audio huddle with your management team, 24/7/365. No links. No calendar lag. Faster and easier than Zoom.

VISCERAL PROOF

2. Wisdom Twin

Role-specific judgment, not a cloned person. Trained on 10-20 years of that role's calls, notes, email, docs, SOPs, and filings: principles, exceptions, thresholds, escalation rules.

CALLABLE JUDGMENT

3. Trust Layer

Cited evidence, source permissions, named human validation, safe decline, audit logs, human override. Private VPC / on-prem / air-gapped.

WHY REGULATORS SAY YES

One press starts the decision. An authorized leader opens a huddle or async review with precedent, policy, and escalation path already in place.

HOW IT WORKS

Ingest → Twin → Operate. Evidence becomes judgment, under human control.

1

Ingest

10-20 years of emails, meetings, docs, SOPs, filings. Immutable evidence objects with source-level permissions.

2

Twin

Decision episodes reconstructed: why, not just what. The role is the unit, not the person. Nothing authoritative until a named human signs off.

3

Operate

One-press huddles and async reviews. Cited answers, policy gates, confidence decomposition, safe decline. Every action audited. Decision → Audit → Smarter Twin.

LIVE DEMO

www.wisdomtwin.ai/demo: 1-4 minute audio huddles. This is the meeting investors take the meeting for.

An RFP is due this week. The Sales Director can't attend.

His twin does.

1

CEO opens a one-press huddle: 60 pages of compliance requirements, no proposal team available.

2

The Sales Director's Wisdom Twin joins instantly: trained on 20 years of proposals, RFPs, compliance calls, and email.

3

The twin, with citations: **"We can do this. I can win this, this week."** Policy-gated, transcribed, auditable.

4

Every interaction: source-cited, permission-aware, safe-decline enabled.

CLIENT OUTCOME

We used our own twin to win a government RFP for Talent Lab. Talent Lab's CEO said on our podcast the win generated \$5.5M in revenue for her company in 2025. A client outcome on the record, not product revenue.

Governance is not a feature. It is the product.

Cited evidence

Every answer traces to authorized sources.

Policy checks

Gated against policy version and approval state.

Named validation

A human expert signs off. Accountability has a name.

Confidence decomposition

The reasoning behind confidence, not one score.

Safe decline

Refuses when judgment is unsupported:
"I will get the human."

Source-level permissions

Access follows the evidence, not the model.

DETERMINISTIC ENFORCEMENT OUTSIDE THE MODEL. The model cannot self-approve a restricted action.

Regulated enterprises need speed, but cannot trade away control.

- **The deployment blocker:** financial services, insurance, legal, healthcare administration, and public-sector teams must preserve provenance, permissions, policy context, auditability, and named accountability.
- **What they can't use:** public copilots that hallucinate, leak data, or bypass approvals.

Copy-pasting judgment into ChatGPT is a compliance violation waiting to happen. Governed speed must be provable to a risk committee.

The Trust Layer makes speed deployable.

WisdomTwin is not generic chat. It is a governed decision-speed platform, built so a regulated buyer can say yes when public AI is off the table.

Private cloud

On-prem

Air-gapped

Model-agnostic

Banking · Insurance · Healthcare · Legal · Government ·
Crown corporations

COMPETITION

Everyone else retrieves. We govern judgment.

Capability	MS Copilot	ServiceNow	Palantir / Glean	WisdomTwin
Decision Episode Schema (why, not what)	NO	NO	PARTIAL	YES
Deterministic policy + permission enforcement outside the model	PROMPT	WORKFLOW	CUSTOM	CODE
Named human validation, safe decline, audit trail	NO	TICKET	HUMAN LOOP	BUILT-IN
Private cloud / on-prem / air-gapped + model-agnostic	CLOUD	CLOUD	LIMITED	YES
Can it attend the meeting instead of a human, with cited judgment?	Summarizes	Routes	Analyzes	ATTENDS + GOVERNS

BUDGET REFRAME

Keyman insurance pays cash after the loss: months of disruption, no continuity. WisdomTwin preserves the judgment from day one, and the asset compounds. Operational hedging, not financial hedging.

MOAT

Every governed decision makes the organization faster.



01 Decision Episode Schema

Easy to copy the schema,
hard to populate it.

02 Role Wisdom Graph

Captures why, not just what.
Compounds per customer.

03 Judgment Policy Engine

Code the model can't talk
past.

04 Judgment Evaluation Dataset

Unbuyable without the
customer. The only asset
that proves fidelity.

A competitor can ship the interface in a quarter. They cannot ship years of validated judgment. No patents claimed, no exclusive data claimed: the moat is the compounding validation dataset plus the governance runtime.

MARKET

Every regulated decision currently bottlenecked by synchronous meetings.

Banking & capital markets

Insurance

Healthcare / Pharma admin

Legal & compliance

Government & crown corporations

Government & military: the extreme bottleneck

Roughly **24% of GDP across the G20**, under **1% of AI token usage** today (internal estimate). Command and procurement decisions need seconds and minutes, not days of meetings. Sovereign buyers cannot use public chatbots: air-gapped is table stakes.

Serviceable market (internal model)

\$6-12B SAM in 2026: 20,000-30,000 mid-market and enterprise firms with 10+ executives (finance, healthcare, pharma, PE portfolio companies) at \$250-400k annual. **Every 200 organizations at the \$500k/yr tier is \$100M ARR.** 20 validations prove the model.

Beachhead: Toronto

Densest concentration of regulated buyers plus AI talent in North America. Founder-led sales, tight loop.

Expansion: MENA via Dubai

EDM Association: 700+ firms, 30k data professionals. UAE government advisory channel. We ARE async management: already operating 24/7 across SF-Toronto and Dubai-Toronto.

CATEGORY MAP

Everyone sells async communication. We sell async decisioning.

SOLD AS SAAS ↑

Async comms tools

Loom, Twist, Yac, GitLab handbook, Slack clips. **They cut meetings. They don't make decisions.**

WISDOMTWIN — alone here

Async decisioning as SaaS: twins trained on 10-20 years of a role's judgment join a 2 AM huddle with audit trails, named validation, human override. Air-gapped deployable.

Creator & industrial twins

Antix avatars; Siemens (9.4% share), NVIDIA, Dassault: factories, cities, supply chains. **Zero leadership decision authority.**

Internal exec-twin experiments

Uber "Dara AI", Toyota "AI Akio", Big-4 personas, Simile (\$200M) x CVS, Viven (\$35M). **Closest to judgment. Locked inside. Not products.**

What the map says

Async comms proved the behavior, then stopped at communication. The serious judgment work is internal-only: it validates demand and is not for sale. Funding already flows to twins of people (\$200M Simile, \$35M Viven), all aimed at research, HR, comms.

Why we hold the quadrant

Governance is the product: evidence, permissions, audit, override. And the customer-specific judgment dataset is unbuyable without the customer.

← ASYNC COMMUNICATION

ASYNC DECISIONING →

A \$8-12B serviceable market today. A \$50B+ category as retirements peak.

2026: the behavior already exists

KLARNA: CEO's digital self presented earnings

ZOOM: CEO avatar opened the earnings call

KPMG · DELOITTE · MCKINSEY: 10,000s of personas

SIMILE x CVS: 400,000 twins live

Public reports, 2025-2026. Proof the behavior exists; none of these own async management.

Bottom-up SAM (internal model)

Mid-market + enterprise firms, 10+ exec roles, regulated bias, NA + EU

20-30k firms

x ~\$300k average ARR (locked ladder: \$250k pilot to \$500k+/yr)

~\$300k

2026 serviceable market, with 2-3x regulated

\$8-12B

2030: the premium shifts to "who makes the call"

Demographics: 10,000 Americans turn 65 daily; 18% of the US is 65+ by 2030 (US Census). Mass retirements drain institutional judgment.

AI acceleration: agents absorb workflows while leadership latency stays 9-to-5. Digital twins ~\$3.9B (2024) to ~\$62B (2034), analyst estimates.

Expansion: from exec teams to ops, legal, and sales leadership. Regulated sectors pay 2-3x for governance.

\$50B+

leadership-decisioning category by 2030 (internal estimate)

\$100M

ARR per 200 organizations at the \$500k/yr tier

One price per stage. One unit: the role.

Stage	Investment	Term	Roles
Judgment Assessment	\$0	Entry	Scoping
Concept Validation	\$50,000	2 months	10
Structured Pilot	\$250,000	6 months	50
Commercial	\$500,000+/yr	Annual	500

No tokens. No meters.

You pay for judgment capacity deployed, never per question asked. Beyond included roles: \$5,000 one-time setup + \$500/month per role. Predictable budget: no usage punishment for adoption.

Ontario reimbursement (pilot de-risk)

COJG/OJAP can reimburse the client up to \$250k per organization where eligible, making net pilot cost \$50-90k. A reimbursement to the client, not a discount on enterprise value. Billing in USD.

The \$1M raise and the first revenue stage are the same number: 20 Concept Validations at \$50k.

TRACTION · HONESTY BUILDS CREDIBILITY

Working demos. Pre-revenue. A clear next milestone.

0

production users

0

paying customers

\$0

product revenue: stated plainly

5

industry demo scenarios,
founder-built end to end

What exists today

Multi-speaker audio huddles (recorded, transcribed) · evidence timelines and decision episodes · human validation flow · role-scoped queries with citations and confidence decomposition · policy denial and safe decline · held-out evaluation framework. Live: www.wisdomtwin.ai/demo

Next milestone

Controlled lighthouse pilots: one organization, one critical role, one authorized evidence set, a named success gate.

CLIENT OUTCOME

We used our own twin to win a government RFP for Talent Lab. Its CEO said on our podcast the win generated \$5.5M in revenue for her company in 2025. Attributed, on the record, and never presented as product revenue.

TEAM · BUILDER + COMPLIANCE ARCHITECT

We ARE asynchronous management before we sell it.

Roman Bodnarchuk

Co-Founder & CEO · SF-Toronto · full-time, writes the code

- 28 years turning executive judgment into systems: founded N5R in 1998; P&G digital agency of record for 6 years
- Enterprise scale-up operator across marketing and technology, 15 countries
- Personally built the current WisdomTwin demos end to end; leads product, fundraising, and founder-led sales

Stella Cabrera

Co-Founder, Governance & Institutional Trust · Dubai-Toronto

- 20+ years with banks, fintechs, and regulators; Chairman & Past President, RMA Toronto; co-author, RMA x OSFI operational resilience guidance
- Senior Advisor & MENA Lead, EDM Association (700+ firms, 30k data professionals); consulting to the UAE government; full-time on close of pre-seed
- "We understand the risk committee because I already chair one."

Roman builds the twin; Stella makes it deployable. Incoming: Deborah Kaplan (former CRO, SAP Sustainability) joins as CFO post-financing. 24/7 coverage across time zones is not a slide claim: it is how we work.

The raise is a sales plan with dates. The seed gate is named in advance.

M0

Close

\$1M SAFE closes. Stella converts to full-time. Governance hardening and source connectors begin.

M1-3

First validations

Evaluation dataset + customer-hosted deployment profiles shipped. First 5 Concept Validations signed in Toronto: founder-led, grant-assisted.

M4-9

Scale validations

15-20 Concept Validations at \$50k. Controlled lighthouse pilots run against named success gates measured on customer baselines.

M10-18

Convert

Validations convert to \$250k Structured Pilots and the first \$500k/yr Commercial cohort. MENA channel opens via Dubai. **Seed-round gate: 3 lighthouse pilots passing held-out fidelity evaluation + \$500k contracted. \$1M = 18 months of founder-led runway.**

SEED GATE: 15-20 validations signed + lighthouse pilots passing held-out fidelity gates + first commercial conversions. When validations convert to \$250k pilots, the seed prices itself. (Forward-looking plan, not a projection of measured results.)

THE ASK

Raising \$1M USD Pre-Seed via SAFE.

The raise and the first revenue stage are the same number. \$1M funds the path to 15-20 Concept Validations at \$50k, converting to \$250k Structured Pilots and the first \$500k/yr Commercial cohort.

Governance controls + audit infrastructure

Source connectors + evaluation dataset

Private VPC / on-prem / air-gapped

Forward-deployed engineering

Founder-led GTM: Toronto + MENA

Preserve institutional judgment before it walks out the door.

Every employee, backed by a Wisdom Twin.

WE SELL SPEED. ONE PRESS MOVES THE DECISION FORWARD.



NO MEETINGS. Your twin goes so you don't have to.